

1. C program to illustrate effect of setjmp and longjmp functions on register, volatile and automatic variables.

```
#include <setjmp.h>
#include<stdio.h>
#include<stdlib.h>

static void f1(int, int, int, int);
static void f2(void);

static jmp_buf jmpbuffer;
static int globval;

int main(void)
{
    int autoval;
    register int regival;
    volatile int volaval;
    static int statval;

    globval = 1; autoval = 2; regival = 3; volaval = 4; statval = 5;
    if (setjmp(jmpbuffer) != 0)
    {
        printf("after longjmp:\n");
        printf("globval = %d, autoval = %d, regival = %d, volaval = %d, statval = %d\n", globval,
autoval, regival, volaval, statval);
        exit(0);
    }
    /*
    * Change variables after setjmp, but before longjmp.
    */
    globval = 95; autoval = 96; regival = 97; volaval = 98;
    statval = 99;
    f1(autoval, regival, volaval, statval); /* never returns */
    exit(0);
}

static void f1(int i, int j, int k, int l)
{
    printf("in f1():\n");
    printf("globval = %d, autoval = %d, regival = %d, volaval = %d, statval = %d\n", globval, i, j,
k, l);
    globval=10000;
    j=10000;
    f2();
}

static void f2(void)
{
    longjmp(jmpbuffer, 1);
}
```

Explanation

This program demonstrates the behavior of `setjmp` and `longjmp` on different variable storage classes:

1. Initialization:

- Global variables:
 - `globval` is initialized to 1.
- Automatic variable:
 - `autoval` is initialized to 2.
- Register variable:
 - `regival` is initialized to 3 (compiler might store it in a register).
- Volatile variable:
 - `volaval` is initialized to 4.
- Static variable:
 - `statval` is initialized to 5.

2. `setjmp(jmpbuffer)`:

- The call to `setjmp(jmpbuffer)` saves the current execution state (including register values) in the `jmpbuffer` buffer.
- It returns 0 on the first call, indicating the initial jump context.

3. Variable Modification:

- **After `setjmp` but before `longjmp`:**
 - The program modifies the values of all variables:
 - `globval` is changed to 95.
 - `autoval` is changed to 96.
 - `regival` is changed to 97.
 - `volaval` is changed to 98.
 - `statval` is changed to 99.

4. `f1(autoval, regival, volaval, statval)`:

- The program calls `f1` with the current values of the arguments (96, 97, 98, 99). However, due to `longjmp` inside `f1`, this function never returns control to `main`.

5. Inside `f1`:

- `f1` prints the values of the arguments (which are the initial values from `main` after `setjmp`).
- `f1` modifies `globval` to 10000.
- `f1` modifies the second argument (`j`, which is a copy of `regival`) to 10000 (this change won't affect the actual `regival` in `main`).

6. `f2()`:

- `f1` calls `f2()`.
- `f2` calls `longjmp(jmpbuffer, 1)`. This does the following:

- Restores the execution state saved in jmpbuffer during the first setjmp call.
- Jumps back to main immediately after the setjmp call, as if f1 and f2 never happened.
- The value of the argument (1) passed to longjmp is ignored in this program.

7. Back in main (after longjmp):

- The program execution resumes in main as if the function calls to f1 and f2 never occurred.
- The modified values of the variables before longjmp are no longer relevant.
- main prints the values of the variables after setjmp:
 - globval will be 95 (modified before longjmp).
 - autoval will be undefined (stack memory overwritten during function call/return).
 - regival will have an unpredictable value (compiler might have used a register).
 - volaval might be 98 (compiler's optimization behavior), but could also be the original value (4) depending on the compiler.
 - statval will be 99 (static variable retains its value).

8. Program Termination:

- main exits with exit(0).
2. Write a C program To create a child process and child should perform addition of two numbers and parent process to perform product of two numbers.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
```

```
int main() {
    int num1 = 5, num2 = 3;

    int pid = fork();

    if (pid < 0) {
        perror("fork");
        exit(1);
    }

    else if (pid == 0) { // Child process
        int sum = num1 + num2;
        printf("Child process: Addition of %d and %d is %d\n", num1, num2, sum);
```

```

        exit(0);
    }
else { // Parent process
    int product = num1 * num2;

    printf("Parent process: Multiplication of %d and %d is %d\n", num1, num2, product);

    wait(NULL); // Wait for child process to finish
}

return 0;
}

```

Explanation:

Initial State:

- Parent process starts execution.
- Variables:
 - num1 = 5
 - num2 = 3

fork() System Call:

1. The parent process calls fork().
2. The kernel creates a child process that is a nearly identical copy of the parent process.
3. Both processes (parent and child) now exist and continue execution from the same point after fork().

Variable Copies:

- Each process (parent and child) gets its own copy of the variables.
 - num1 (value 5) and num2 (value 3) are copied into both processes' memory spaces.

Process Divergence:

1. The parent process checks the return value of fork():
 - If the return value is less than 0 (pid < 0), an error occurred during fork(). The program exits with an error message.
 - If the return value is 0 (pid == 0), it's the child process.
 - If the return value is greater than 0 (pid > 0), it's the parent process with the child's process ID (pid) stored in the variable pid.
2. **Child Process (pid == 0):**
 - The child process continues execution.
 - It calculates the sum of num1 and num2: $\text{sum} = \text{num1} + \text{num2} = 5 + 3 = 8$.
 - The child process prints the calculated sum:
 - Child process: Addition of 5 and 3 is 8

- The child process exits successfully (exit(0))
- 3. **Parent Process (pid > 0):**
 - The parent process continues execution after the child process exits.
 - It calculates the product of num1 and num2: $\text{product} = \text{num1} * \text{num2} = 5 * 3 = 15$.
 - The parent process prints the calculated product:
 - Parent process: Multiplication of 5 and 3 is 15
 - The parent process waits for the child process to finish using wait(NULL).
 - The parent process exits successfully (return 0).

Program Termination:

- Both child and parent processes terminate, and the program execution ends.

Output:

Child process: Addition of 5 and 3 is 8

Parent process: Multiplication of 5 and 3 is 15

3. Write a C program To create a child process and both parent and child should share a text file to perform read and write operations on a file. Show how the offset value is shared between parent and the child processes.

```
#include <stdio.h>

#include <stdlib.h>

#include <unistd.h>

#include <fcntl.h>

#include <sys/mman.h>

#include <sys/stat.h>

#define FILE_NAME "shared_file.txt"

#define SHARED_MEM_SIZE sizeof(int)

int main() {

    int fd;

    int *offset_ptr;

    char *write_buffer = "This is written from the parent process.\n";

    char read_buffer[100];

    // Create/open the file with read/write permissions

    fd = open(FILE_NAME, O_RDWR | O_CREAT, 0666);
```

```

if (fd == -1) {
    perror("open");
    exit(1);
}

// Create a shared memory segment
int shm_fd = shm_open("offset_shm", O_RDWR | O_CREAT, 0666);
if (shm_fd == -1) {
    perror("shm_open");
    exit(1);
}

// Set the size of the shared memory segment
ftruncate(shm_fd, SHARED_MEM_SIZE);

// Map the shared memory segment to the address space of this process
offset_ptr = mmap(NULL, SHARED_MEM_SIZE, PROT_READ | PROT_WRITE,
MAP_SHARED, shm_fd, 0);
if (offset_ptr == MAP_FAILED) {
    perror("mmap");
    exit(1);
}

pid_t pid = fork();

if (pid < 0) {
    perror("fork");
    exit(1);
} else if (pid == 0) { // Child process
    printf("Child process started.\n");

    // Seek to the current offset in the file

```

```

lseek(fd, *offset_ptr, SEEK_SET);

// Read from the file
int bytes_read = read(fd, read_buffer, sizeof(read_buffer));
if (bytes_read > 0) {
    read_buffer[bytes_read] = '\0'; // Null-terminate the string
    printf("Child process: Read %d bytes: %s\n", bytes_read, read_buffer);
} else {
    perror("read");
}

// Update the offset with the current position in the file
*offset_ptr = lseek(fd, 0, SEEK_CUR);

printf("Child process finished.\n");
} else { // Parent process
    printf("Parent process started.\n");

    // Write to the file
    write(fd, write_buffer, strlen(write_buffer));

    // Update the offset with the current position in the file
    *offset_ptr = lseek(fd, 0, SEEK_CUR);

    printf("Parent process finished.\n");
}

// Unmap the shared memory segment
munmap(offset_ptr, SHARED_MEM_SIZE);

// Close the shared memory segment descriptor

```

```

close(shm_fd);

// Close the file descriptor
close(fd);

return 0;
}

```

Explanation

Initial State:

- Parent process starts execution.
- Shared memory segment "offset_shm" is created with size sizeof(int).
- File shared_file.txt is opened/created with read/write permissions.
- offset_ptr points to the mapped shared memory segment containing the offset value (initially 0).

fork() System Call:

1. The parent process calls fork().
2. The kernel creates a child process, which is a nearly identical copy of the parent.
 - Both processes (parent and child) now have their own copy of variables but share the same file descriptor (fd) and mapped shared memory segment (offset_ptr).

Process Divergence:

1. **Parent Process (pid > 0):**
 - Prints "Parent process started."
 - **Write Operation:**
 - Writes the content "This is written from the parent process.\n" to the file using write.
 - Updates the offset value in the shared memory segment (*offset_ptr) with the current position in the file using lseek(fd, 0, SEEK_CUR). This points to the end of the written data.
 - Prints "Parent process finished."
2. **Child Process (pid == 0):**
 - Prints "Child process started."
 - **Read Operation:**
 - Seeks to the current offset value in the file using lseek(fd, *offset_ptr, SEEK_SET).
 - **Scenario 1 (Child Reads After Parent Writes):**
 - If the parent has already written, the child will read the written content (This is written from the parent process.\n).
 - **Scenario 2 (Child Reads Before Parent Writes):**
 - If the child reads before the parent writes, the read operation will return 0 bytes (read_buffer will be empty).
 - Prints the read data (if any) along with the number of bytes read.

- Updates the offset value in the shared memory segment (*offset_ptr) with the current position in the file (which might be the same as before or at the end of the read data).
- Prints "Child process finished."

Synchronization with Shared Memory:

- The offset value in the shared memory segment acts as a coordination point between the parent and child processes.
- The parent writes, updates the offset, and then the child reads from that updated position.
- This ensures that the child doesn't overwrite the parent's written data and can potentially read the written content (depending on the timing).

Cleanup:

- Both processes unmap the shared memory segment, close the shared memory segment descriptor, and close the file descriptor before exiting.

Output:

The output will vary depending on the timing of the read operation in the child process, as explained in the previous explanation